

# Telemedicine Scheduling & Virtual Visit Platform

## Submitted By:

**Name:** Bhumi Dwivedi

**Roll No.:** 44

**Registration No.:** 12407101

**Course:** B.Tech – Generative AI

**Semester:** 3

**Session:** 2024–2025

**College:** Lovely Professional University

---

## Overview

### Telemedicine Scheduling & Virtual Visit Platform

Enables patients to book appointments, conduct virtual video visits, and receive associated documentation (prescriptions, notes, post-visit summaries).

### Primary Objectives

- High appointment integrity (accurate availability, no double booking)
  - Fast video join performance
  - Maintainable system boundaries (Scheduling, Video, Records/Documentation)
- 

## Stakeholder Analysis & Prioritization

### Stakeholders

Stakeholder	Interests	Priority
Patients	Easy booking, reliable video join, secure data	High
Doctors	Efficient scheduling, stable video, record access	High
Admins	Configuration, compliance, monitoring	Medium

<b>Compliance/Legal</b>	Privacy, audit trails, data residency	High
<b>Support Ops</b>	Troubleshooting and monitoring	Medium

## Key Stakeholder Needs

- Patients: availability clarity, low-friction workflows
  - Doctors: predictable schedules, minimal tech friction
  - Compliance: auditability, retention policies
  - Admins: updates, overrides, clinic setup
- 

# Functional Requirements

## Core Functionalities

### 1. Appointment Discovery

- Search slots by doctor, specialty, date, insurance.

### 2. Booking

- Reserve → Confirm → Notify workflow.
- Cancellation & rescheduling.

### 3. Pre-Visit

- Patient reminders.
- Doctor dashboard view of schedule.

### 4. Virtual Visit

- Tokenized video link generation (third-party provider).
- 1-click join experience for patients and doctors.

### 5. Post-Visit

- Prescriptions and clinical notes.

- Automated visit summary to patient.

## 6. Admin Functions

- Schedule templates, availability overrides.
  - Reporting & auditing.
- 

# Non-Functional Requirements

## Availability

- 99.9% uptime during clinic operating hours.

## Performance

- p95 video join time < **3 seconds**.
- p95 slot search < **300 ms**.

## Security & Compliance

- HIPAA / GDPR aligned.
- Data residency per region.
- End-to-end encrypted video sessions.

## Reliability & Observability

- End-to-end audit logs (access, actions).
- Structured logs, metrics (RED/USE), distributed tracing.

## Scalability

- Burst handling during peak morning hours (5–10× load).
-

# Constraints & Assumptions

## Constraints

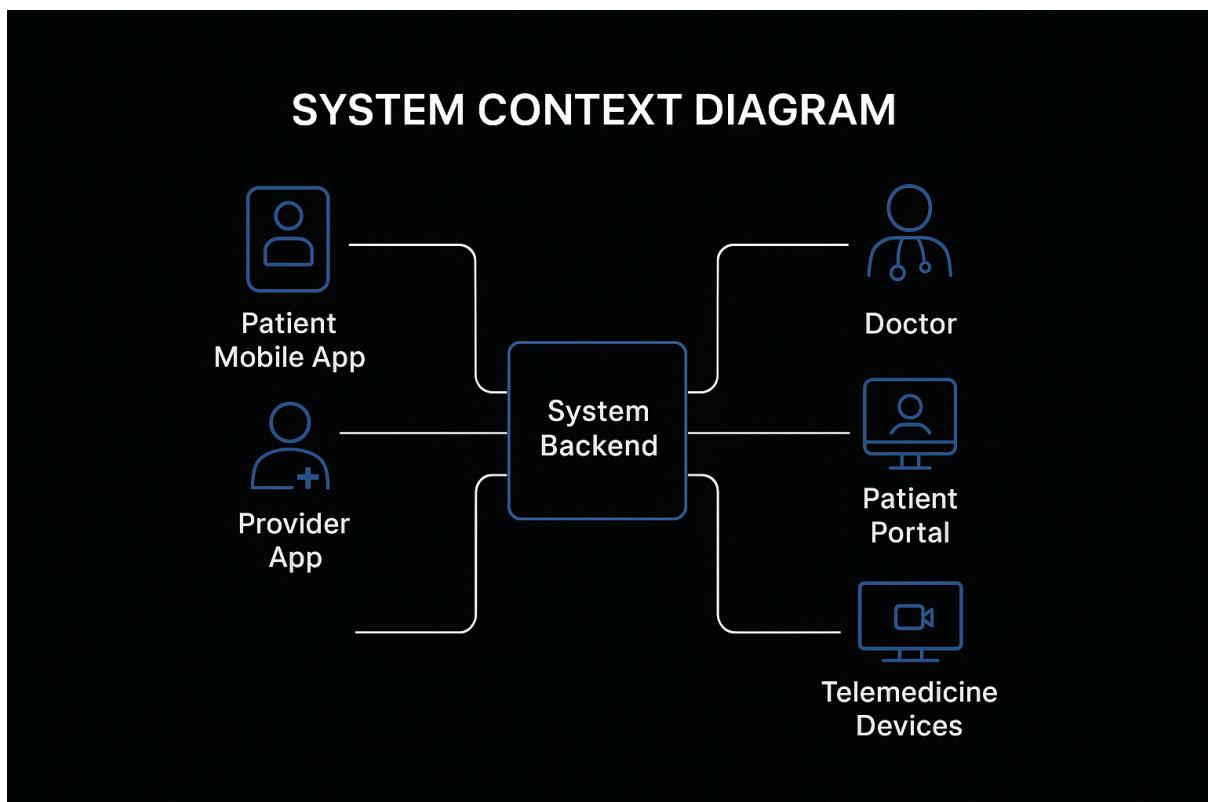
- Third-party video API must be used.
- Regional data stores — no cross-region PII replication unless anonymized.
- Clinic hours vary; platform must support arbitrary schedules.

## Assumptions

- Stable third-party video provider with SLA  $\geq 99.9\%$ .
- Users have stable internet connections.
- No in-house EHR; platform integrates with external EHR over standard APIs.

---

## System Context Diagram



**Actors:** Patient, Doctor, Admin, Compliance

## Systems

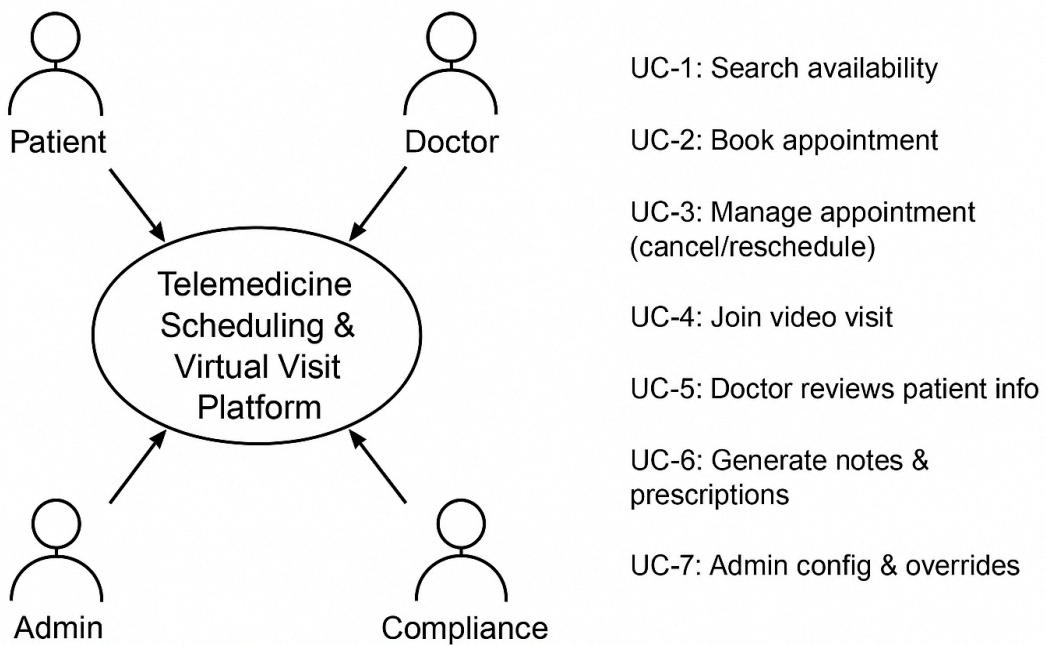
- **Telemedicine Platform**
  - Scheduling Service
  - Video Gateway Service
  - Records/Notes Service
  - Notification Service
  - Identity & Access
- **External Systems**
  - EHR (medical history, visit notes sync)
  - Third-party Video Provider
  - Notification Gateway (SMS/email)
  - Payment (optional)"

## Flows

- User ↔ Identity
- Platform ↔ EHR (sync notes/prescriptions)
- Platform ↔ Video API (create tokens)
- Platform ↔ Notification API (SMS/email)

---

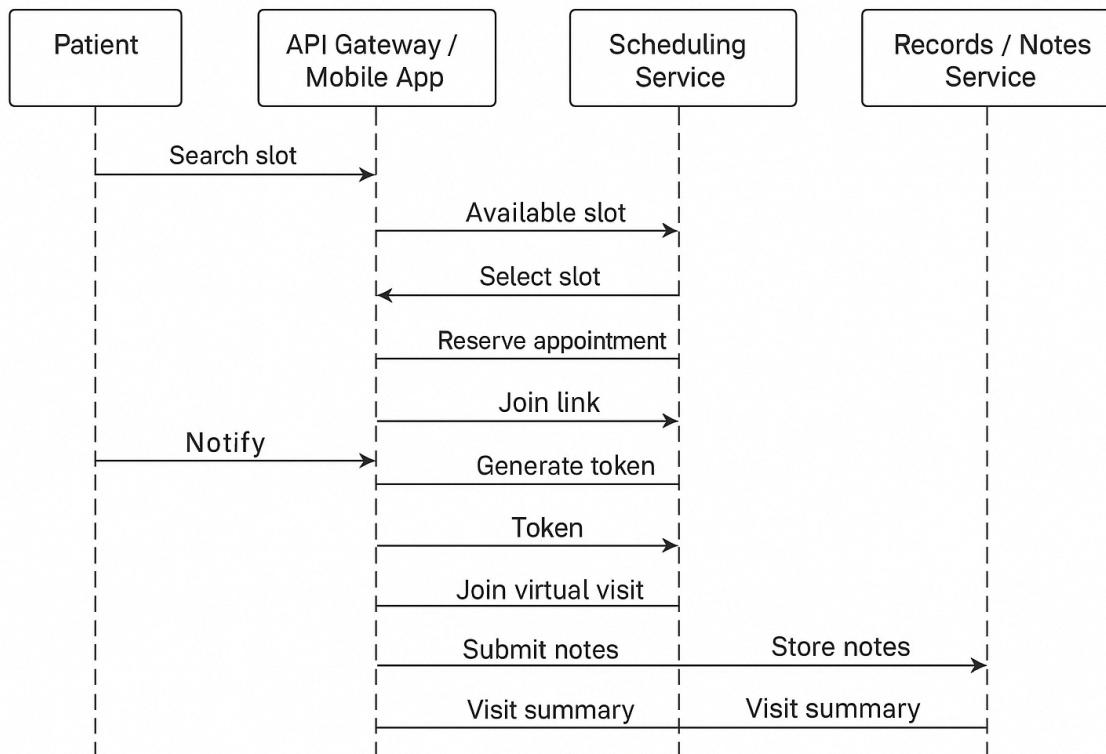
## Use Case Diagram + User Stories



## User Stories

1. As a patient, I want to search for the earliest available slot so that I can book quickly.
2. As a patient, I want a simple link to join my video visit.
3. As a doctor, I want to view my daily schedule with patient details.
4. As an admin, I want to modify availability templates.
5. As compliance, I need audit logs for each access and event.

## Sequence Diagrams



## Booking Sequence

1. User searches availability → Scheduling Service queries Slot Store + Cache.
2. User selects slot → Scheduler places a **temporary hold** (TTL locking).
3. Payment (optional) / confirmation → Appointment persisted.
4. Notification Service sends SMS/email.

## Video Join

1. User clicks join link → Identity verifies token.
2. Video Gateway Service requests session token from provider.
3. Returns scoped join token to client.
4. Client establishes WebRTC/video session.

## Post-Visit Notes

Doctor submits notes → Records Service stores → Sync with EHR → Summary sent to patient.

---

# High-Level Architecture

Patient / Doctor / Admin ---> API Gateway ---> Scheduling ---> Video ---> Records ---> Notifications ---> Payments

## Services

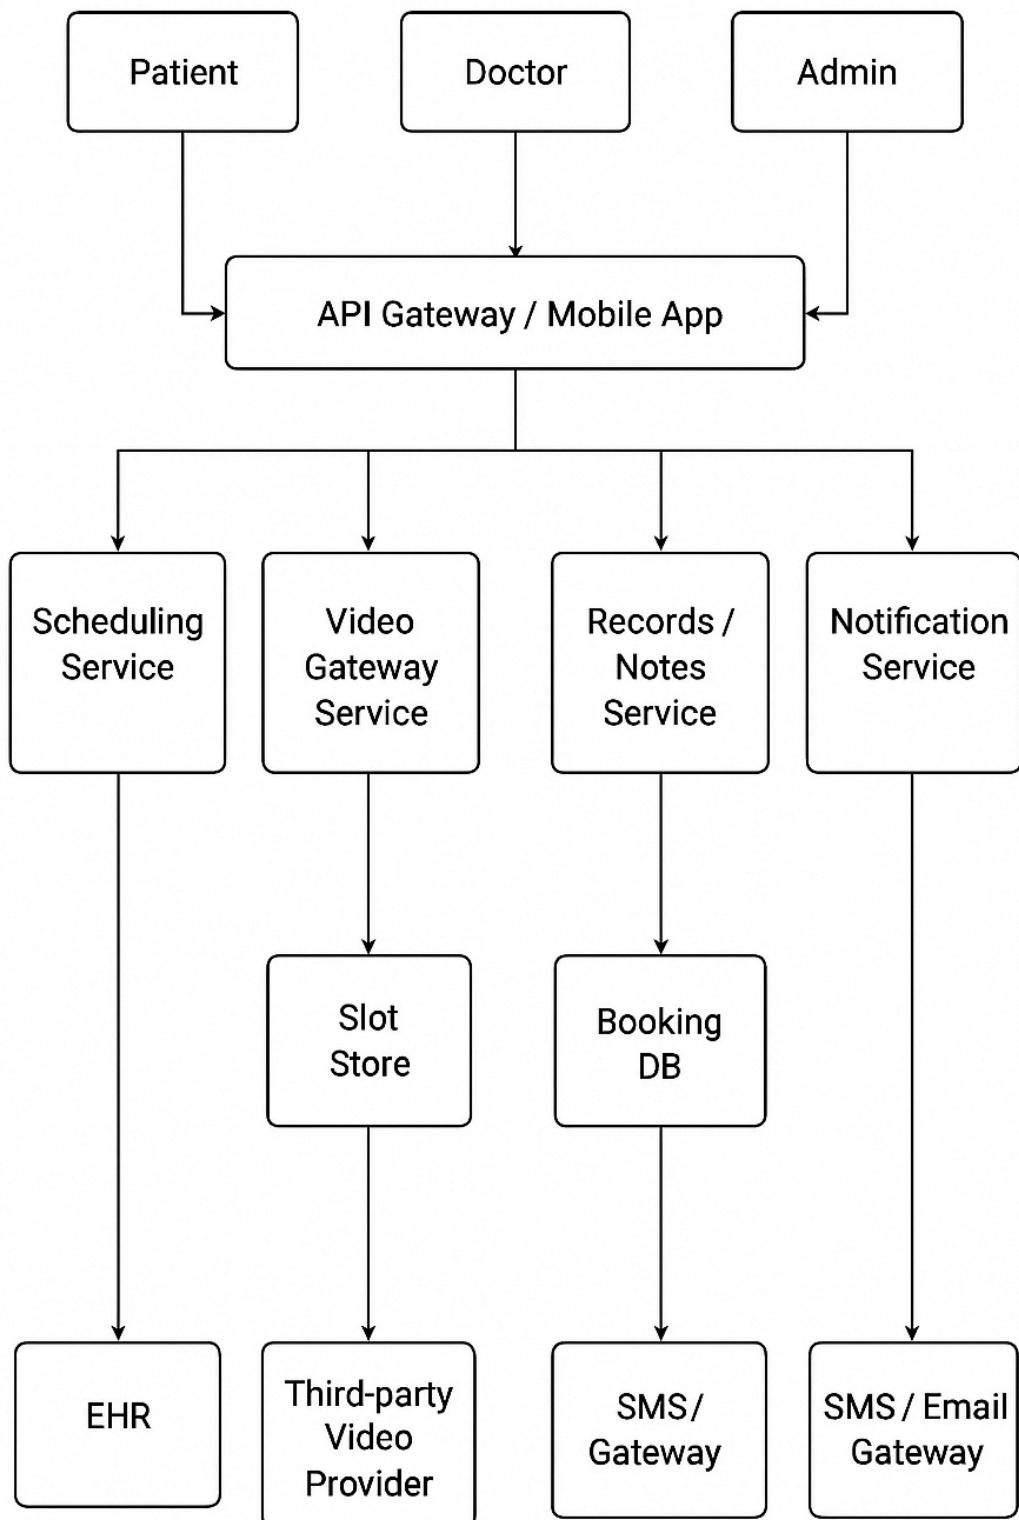
- **Scheduling Service:** Slot management, booking, holds.
- **Video Gateway Service:** Abstract video vendor (Twilio/Zoom/etc).
- **Records Service:** Notes, prescriptions, visit summaries.
- **Notification Service:** Email/SMS push.
- **Identity Service:** OAuth2/OIDC, MFA optional.
- **Admin Service:** Templates, overrides.

## Data Stores

- **Booking DB** (SQL, partitioned by clinic/region)
- **Slot Store** (Redis or DynamoDB for quick reads)
- **Records DB** (document store: notes, attachments)
- **Audit Log Store** (append-only)

## Infra

- API Gateway
- CDN for static app content
- Event Bus (Kafka/SQS) for async tasks
- Cache layer (Redis)
- Observability stack (Prometheus + Loki + Grafana)



---

# Engineering Notes

## Capacity Planning

- Morning peak: assume 10k concurrent users for mid-sized provider.
- Slot search QPS: ~300–1000/s.
- Video join bursts: ~200–400/s.

## Back-of-the-Envelope

- Redis caches ~2–5 MB/day of slot metadata.
- Booking DB read-heavy (10:1 read/write).
- GP2/PostgreSQL or Aurora recommended.

## API Specs (Examples)

**GET /v1/slots?doctorId=&date=**

**POST /v1/appointments**

**POST /v1/video/join-token**

**POST /v1/notes**

## Data Model

**Slot:** id, doctorId, start, end, status (open/held/booked).

**Appointment:** id, slotId, patientId, status.

**Visit:** id, appointmentId, videoSessionId.

**Note:** visitId, authorId, content, attachments.

## Consistency Choices

- **Strong consistency** on booking transactions.
- **Eventual consistency** for summaries, notifications.

## Caching & Indexing

- Slot search in Redis.

- DB indexes: (doctorId, start\_time), (patientId, date).

## Rate Limiting

- Per-user & per-IP throttles at API gateway.
- Burst protection for join-token generation.

## Resiliency

- Retries w/ exponential backoff.
- Timeouts on video provider > 2s.
- Circuit breakers on external APIs.

## Observability

- Traces from frontend → backend → video provider.
- Alerts: booking error spike, join-token latency, slot cache hits < 95%.

## Maintenance

- Zero-downtime deploys w/ blue-green.
- Daily database integrity checks.
- Rotation of signing keys every 90 days.

# Quality Targets, SLOs, Scalability, Trade-Offs

## Explicit SLOs

Area	SLO
Availability	$\geq 99.9\%$ clinic hours
Slot search latency	p95 < 300 ms
Video join token	p95 < 3 s

Booking atomicity failures < **0.01%**

## Scalability Plan

- Read-heavy operations → CDN + cache.
- Partition database by region/clinic.
- Horizontally scale Scheduling and Video Gateway services.
- Queue-based asynchronous tasks for notifications + record sync.

## Trade-Off Discussion

- **SQL vs NoSQL:** SQL chosen for strong consistency in booking; NoSQL for notes.
- **Third-party video vs in-house:** Outsourced video reduces complexity but limits control; a gateway service helps abstraction.
- **Regional data residency** complicates analytics; solved with anonymization and regional partitions.
- **High availability** increases cost due to multi-AZ setups, but required for medical workflows